

CROSS VALIDATION

MACHINE LEARNING

CROSS-VALIDATION

Cross-Validation (CV) is a **model validation technique** used to check how well your machine learning model will generalize on **unseen data**.

Instead of training on one fixed train-test split, CV **splits the data multiple times** and checks performance in each split.

Purpose

- 1. Reduce overfitting**
- 2. Better model evaluation**
- 3. Reliable performance score**
- 4. Helps in hyperparameter tuning**

Explanation

- Ensures model performs well on unseen data.
- Uses multiple splits instead of one.
- Gives average accuracy/MAE/R² etc.
- Used in GridSearchCV, RandomizedSearchCV.

WHY CROSS-VALIDATION IS USED?

Because a single train-test split is not reliable.

Example:

Train on 80%

Test on 20%

If the 20% test set is **too easy or too difficult**, accuracy becomes misleading.

Cross-Validation solves this by:

- Splitting the dataset into k equal parts (folds)
- Training multiple times
- Averaging the score

This gives **stable and unbiased** performance.

TYPES OF CROSS-VALIDATION

10 are the standard techniques used in ML:

1. K-Fold
2. Stratified K-Fold
3. LOOCV
4. Leave-P-Out
5. ShuffleSplit
6. Repeated K-Fold
7. TimeSeriesSplit
8. Group K-Fold
9. Stratified Group K-Fold
10. Nested Cross-Validation


CROSS-VALIDATION IS USEFUL IN ANY DOMAIN WHERE YOU TRAIN MODELS AND NEED TO EVALUATE THEM, INCLUDING:

AI Domain	Used?	Why
Machine Learning	YES	Evaluate models reliably
Deep Learning	Sometimes	Usually replaced by train/val/test splits due to huge data
NLP	Optional	Depends on dataset size
Computer Vision	Optional	Big datasets → single split is enough
Time Series Forecasting (AI)	YES	Uses TimeSeriesSplit
Reinforcement Learning	Rare	RL uses episodic evaluation
Statistics & Data Science	YES	Model reliability

HOLD OUT CROSS VALIDATION

- The Holdout Method is a **fundamental validation technique** in machine learning used to evaluate the performance of a predictive model.
- The dataset is commonly divided into training set and test set.
- Typical **split ratios** include **70:30, 80:20 or 60:40** depending on dataset size.
- It is **most effective** when the **dataset is large enough** to allow meaningful splitting.
- **Random shuffling** before splitting is often applied to reduce bias.
- Simple and Faster

DRAWBACK: single random split, Not ideal for small datasets



CODE:


```
from sklearn.model_selection import train_test_split  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,  
random_state=42)  
model = LogisticRegression()  
model.fit(X_train, y_train)
```

SHUFFLE SPLIT CROSS VALIDATION

- ShuffleSplit **randomly shuffles** the dataset.
- Then it splits it into **train** (e.g., **70%**) and **test** (e.g., **30%**).
- It does this **multiple times** (e.g., 10 iterations).
- Every iteration gives a **different random 70-30 split**.
(eg: 65 – 35, 80-20, 73-7 ..)
- The model is trained and tested on each different split.
- Finally, we take the **average performance** over all splits.

Benefits: works for small dataset, fast and stable model accuracy

Drawback: time consuming




CODE:

```
from sklearn.model_selection import ShuffleSplit
from sklearn.linear_model import LogisticRegression
#dataset
X, y = make_classification(n_samples=100, n_features=4,
n_informative=2,
n_redundant=0, n_repeated=0, n_classes=2,
random_state=42)
ss = ShuffleSplit(n_splits=5, test_size=0.3, random_state=0)
model = LogisticRegression(random_state=42)
scores = []
for train_index, test_index in ss.split(X):
    X_train, X_test = X[train_index], X[test_index]
    y_train, y_test = y[train_index], y[test_index]
    model.fit(X_train, y_train)
```

LOOCV (LEAVE ONE OUT CROSS VALIDATION)

- In this method the model is **trained on the entire dataset except for one data point which is used for testing**. This process is repeated for each data point in the dataset.
- All data points used for training – low bias
- **very time-consuming for large datasets** as it requires one iteration per data point.

DRAWBACK: Computationally expensive for large datasets.




CODE:

```
from sklearn.model_selection import LeaveOneOut  
# sample dataset  
X = np.array([[1], [2], [3], [4], [5], [6], [7], [8], [9], [10]])  
y = np.array([2, 4, 5, 4, 6, 7, 8, 9, 10, 12])  
# Initialize the LeaveOneOut cross-validator  
loo = LeaveOneOut()  
model = LinearRegression()  
# Perform LOOCV  
for train_index, test_index in loo.split(dataset):  
    # Split data into training and testing sets for the current fold  
    X_train, X_test = X[train_index], X[test_index]  
    y_train, y_test = y[train_index], y[test_index]  
    # Fit the model on the training data  
    model.fit(X_train, y_train)
```

K-FOLD CROSS-VALIDATION

- The dataset is **divided into k equal-sized folds**.
- The model is **trained on k-1 folds and tested on the remaining fold**.
- This process is **repeated k times**, with each fold serving as the test set exactly once.
- The **results are then averaged**.
- **Best use case:** Small to medium datasets
- **Drawback:** slower for large dataset, can't handle imbalanced dataset



CODE:



```
from sklearn.model_selection import KFold, cross_val_score
from sklearn.ensemble import RandomForestClassifier
kf = KFold(n_splits=5, shuffle=True, random_state=42)
model = RandomForestClassifier(random_state=42)
for fold_num, (train_index, test_index) in enumerate(kf.split(X, y)):
    X_train, X_test = X[train_index], X[test_index]
    y_train, y_test = y[train_index], y[test_index]
    model.fit(X_train, y_train)
```

STRATIFIED K-FOLD CROSS-VALIDATION

- Similar to k-fold, but ensures that **each fold maintains the same proportion of target variable classes** as the original dataset.
- This is crucial for imbalanced datasets.
- Best for **imbalanced classification problems**
- Key benefits: **Reliable Metrics, Generalization**

CODE:

```
from sklearn.model_selection import StratifiedKFold
from sklearn.linear_model import LogisticRegression
skf = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
model = LogisticRegression(random_state=42)
for fold_idx, (train_index, test_index) in enumerate(skf.split(X, y)):
    X_train, X_test = X[train_index], X[test_index]
    y_train, y_test = y[train_index], y[test_index]
    model.fit(X_train, y_train)
```




TIME SERIES CROSS-VALIDATION (ROLLING CROSS-VALIDATION)

- Used for **time-dependent data**. Instead of random splits, it respects **the order of data**, using past events to predict future events.
- It's essential for time series forecasting.
- drawback: Can't use future data to predict the past

CODE:

```
from sklearn.model_selection import TimeSeriesSplit
tscv = TimeSeriesSplit(n_splits=5)
for fold, (train_index, test_index) in enumerate(tscv.split(X)):
    X_train, X_test = X.iloc[train_index], X.iloc[test_index]
    y_train, y_test = y.iloc[train_index], y.iloc[test_index]
    model = LinearRegression()
    model.fit(X_train, y_train)
```



|